



Memory Managment - Part II

Amir H. Payberah
payberah@kth.se
Nov. 23, 2022



Reminder



Reminder (1/3)

- ▶ External fragmentation vs. internal fragmentation



Reminder (1/3)


- ▶ External fragmentation vs. internal fragmentation
- ▶ **Compaction**: shuffle memory contents to place all free memory together in one large block.




Reminder (1/3)

- ▶ External fragmentation vs. internal fragmentation
- ▶ **Compaction**: shuffle memory contents to place all free memory together in one large block.
- ▶ Other solutions:
 - Segmentation
 - Paging

Reminder (2/3)

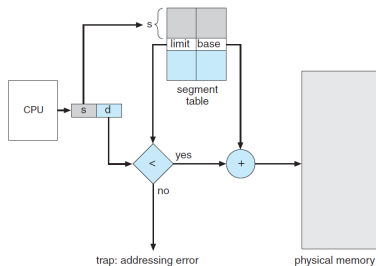


user space


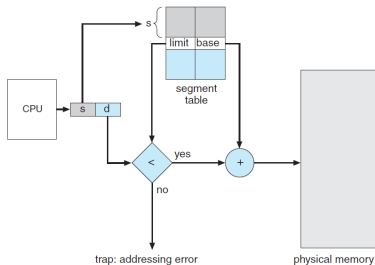


physical memory space

Reminder (3/3)




Reminder (3/3)



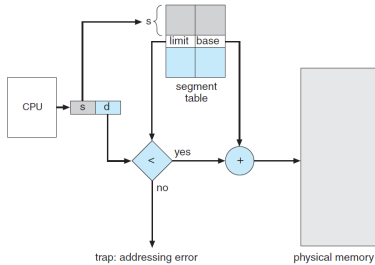


	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table

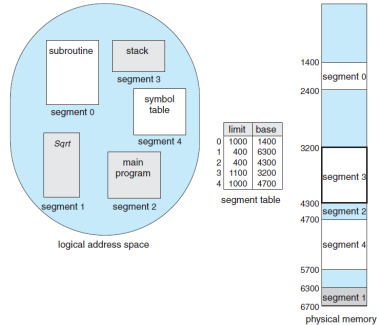



Reminder (3/3)





- A reference to byte 53 of segment 2:

Reminder (3/3)





- A reference to byte 53 of segment 2: $4300 + 53 = 4353$

Reminder (3/3)





- ▶ A reference to byte 53 of segment 2: $4300 + 53 = 4353$
- ▶ A reference to byte 852 of segment 3:

Reminder (3/3)



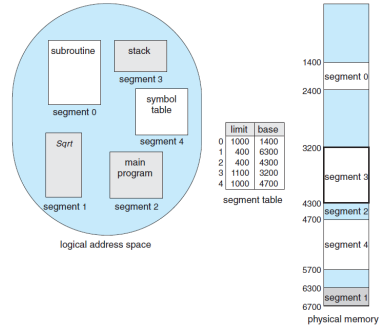
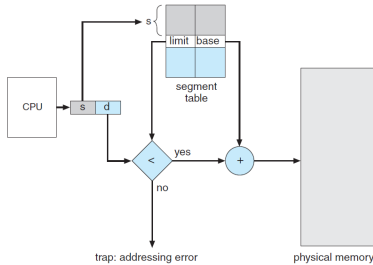
- ▶ A reference to byte 53 of segment 2: $4300 + 53 = 4353$
- ▶ A reference to byte 852 of segment 3: $3200 + 852 = 4052$

Reminder (3/3)



- ▶ A reference to byte 53 of segment 2: $4300 + 53 = 4353$
- ▶ A reference to byte 852 of segment 3: $3200 + 852 = 4052$
- ▶ A reference to byte 1222 of segment 0:

Reminder (3/3)



- ▶ A reference to byte 53 of segment 2: $4300 + 53 = 4353$
- ▶ A reference to byte 852 of segment 3: $3200 + 852 = 4052$
- ▶ A reference to byte 1222 of segment 0: **trap to OS**

Paging



Paging vs. Segmentation

- ▶ **Segmentation** and **paging**, both, permit the physical address space of a process to be **noncontiguous**.



Paging vs. Segmentation

- ▶ Segmentation and paging, both, permit the physical address space of a process to be noncontiguous.
- ▶ Paging avoids external fragmentation and the need for compaction, whereas segmentation does not.



Paging (1/2)

- ▶ Physical address space of a process can be **noncontiguous**.



Paging (1/2)

- ▶ Physical address space of a process can be **noncontiguous**.
- ▶ Divide **physical memory** into **fixed-sized blocks** called **frames**.



Paging (1/2)

- ▶ Physical address space of a process can be noncontiguous.
- ▶ Divide physical memory into fixed-sized blocks called frames.
 - Size is power of 2, between 512 bytes and 16 Mbytes.



Paging (1/2)

- ▶ Physical address space of a process can be **noncontiguous**.
- ▶ Divide **physical memory** into **fixed-sized blocks** called **frames**.
 - Size is **power of 2**, between 512 bytes and 16 Mbytes.
- ▶ Divide **logical memory** into **blocks of same size** called **pages**.



Paging (2/2)

- ▶ Keep track of all free frames.



Paging (2/2)


- ▶ Keep track of all free frames.
- ▶ To run a program of size N pages, need to find N free frames and load program.




Paging (2/2)

- ▶ Keep track of all free frames.
- ▶ To run a program of size N pages, need to find N free frames and load program.
- ▶ Set up a page table to translate logical to physical addresses.


Free Frames



Free Frames



before allocation



after allocation



Address Translation Scheme

- ▶ Logical address generated by CPU is divided into two parts:




Address Translation Scheme

- ▶ Logical address generated by CPU is divided into two parts:
- ▶ Page number (p): used as an index into a page table that contains base address of each page in physical memory.

Address Translation Scheme

- ▶ Logical address generated by CPU is divided into two parts:
- ▶ Page number (p): used as an index into a page table that contains base address of each page in physical memory.
- ▶ Page offset (d): combined with base address to define the physical memory address.




Address Translation Scheme

- ▶ Logical address generated by CPU is divided into two parts:
- ▶ Page number (p): used as an index into a page table that contains base address of each page in physical memory.
- ▶ Page offset (d): combined with base address to define the physical memory address.



- ▶ For given logical address space 2^m and page size 2^n .

Paging Example



Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2


page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory


- The logical address:

Paging Example



- The logical address: $m = 4$ and $n = 2$

Paging Example



- The logical address: $m = 4$ and $n = 2$
- Logical address 3:

Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

- The logical address: $m = 4$ and $n = 2$
- Logical address 3: $5 \times 4 + 3 = 23$

Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2


page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

- The logical address: $m = 4$ and $n = 2$
- Logical address 3: $5 \times 4 + 3 = 23$
- Logical address 10:

Paging Example



- ▶ The logical address: $m = 4$ and $n = 2$
- ▶ Logical address 3: $5 \times 4 + 3 = 23$
- ▶ Logical address 10: $1 \times 4 + 2 = 6$



Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.



Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.
- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes



Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.
- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$



Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.
- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$



Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.
- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$
- ▶ Worst case fragmentation = 1 frame - 1 byte

Paging Example - Internal Fragmentation

- ▶ Still have internal fragmentation.
- ▶ Page size = 2048 bytes
- ▶ Process size = 72766 bytes
- ▶ $\frac{72766}{2048} = 35 \text{ pages} + 1086 \text{ bytes}$
- ▶ Internal fragmentation: $2048 - 1086 = 962 \text{ bytes}$
- ▶ Worst case fragmentation = 1 frame - 1 byte
- ▶ On average fragmentation = $\frac{1}{2}$ frame size



Small Page Size vs. Big Page Size



Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.



Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.



Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.
- ▶ Disk I/O is more efficient when the amount data being transferred is larger (e.g., big pages).

Small Page Size vs. Big Page Size

- ▶ On average fragmentation = $\frac{1}{2}$ page size, hence, small page sizes are desirable.
- ▶ Small pages, more overhead is in the page-table, this overhead is reduced as the size of the pages increases.
- ▶ Disk I/O is more efficient when the amount data being transferred is larger (e.g., big pages).
- ▶ Pages typically are between 4 KB and 8 KB in size.

```
getconf PAGESIZE
```

Page Table Implementation



Page Table

- ▶ Page table is kept in main memory.



Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.



Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.
- ▶ Page-table length register (PTLR) indicates size of the page table.

Page Table

- ▶ Page table is kept in main memory.
- ▶ Page-table base register (PTBR) points to the page table.
- ▶ Page-table length register (PTLR) indicates size of the page table.
- ▶ In this scheme every data/instruction access requires two memory accesses.
 - One for the page table and one for the data/instruction.



Translation Look-aside Buffers (1/2)

- ▶ The two memory access problem can be solved by the use of a special fast-lookup hardware cache called translation look-aside buffers (TLBs).

Translation Look-aside Buffers (2/2)

► TLB

Page #	Frame #

Translation Look-aside Buffers (2/2)

- ▶ TLB

Page #	Frame #

- ▶ Address translation (p, d)

Translation Look-aside Buffers (2/2)

► TLB

Page #	Frame #

- Address translation (p, d)
 - If p is in TLB, get $frame\#$ out.

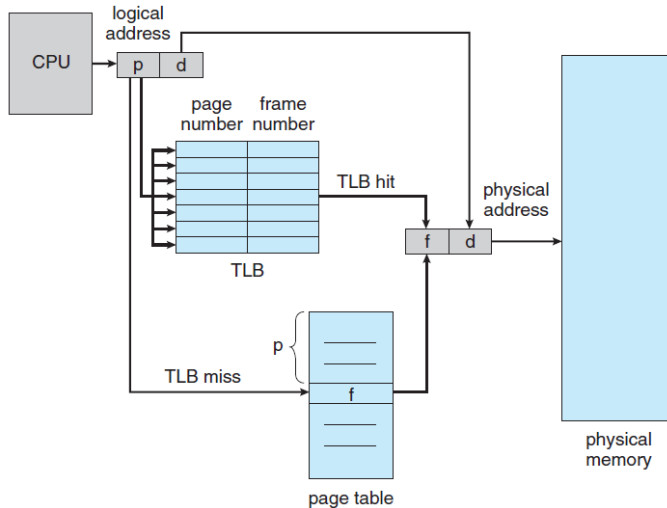
Translation Look-aside Buffers (2/2)

► TLB

Page #	Frame #

- Address translation (p, d)
 - If p is in TLB, get $frame\#$ out.
 - Otherwise, get $frame\#$ from page table.

Paging Hardware With TLB



Structure of the Page Table



Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):



Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):
 - Page size of $4KB = 2^{12}$ ($n = 12$).



Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):
 - Page size of $4KB = 2^{12}$ ($n = 12$).
 - Page table would have 1 million entries $m - n = 20$.

Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):
 - Page size of $4KB = 2^{12}$ ($n = 12$).
 - Page table would have 1 million entries $m - n = 20$.
 - If each entry is $4B$: $4MB$ of physical address space memory for page table alone.

Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):
 - Page size of $4KB = 2^{12}$ ($n = 12$).
 - Page table would have 1 million entries $m - n = 20$.
 - If each entry is $4B$: $4MB$ of physical address space memory for page table alone.
 - That amount of memory used to cost a lot.

Structure of the Page Table (1/2)

- ▶ Consider a 32-bit logical address space ($m = 32$):
 - Page size of $4KB = 2^{12}$ ($n = 12$).
 - Page table would have 1 million entries $m - n = 20$.
 - If each entry is $4B$: $4MB$ of physical address space memory for page table alone.
 - That amount of memory used to cost a lot.
 - Don't want to allocate that contiguously in main memory.



Structure of the Page Table (2/2)

- ▶ Hashed Page Tables
- ▶ Hierarchical Paging

Hashed Page Tables



Hashed Page Tables

- ▶ The logical page number is hashed into a page table.



Hashed Page Tables

- ▶ The logical page number is hashed into a page table.
- ▶ This page table contains a chain of elements hashing to the same location.



Hashed Page Tables

- ▶ The logical page number is hashed into a page table.
- ▶ This page table contains a chain of elements hashing to the same location.
- ▶ Each element contains



Hashed Page Tables

- ▶ The logical page number is hashed into a page table.
- ▶ This page table contains a chain of elements hashing to the same location.
- ▶ Each element contains
 1. The logical page number



Hashed Page Tables

- ▶ The **logical page number** is **hashed** into a **page table**.
- ▶ This **page table** contains a **chain** of elements **hashing to the same location**.
- ▶ Each element contains
 1. The **logical page number**
 2. The value of the **mapped page frame**



Hashed Page Tables

- ▶ The **logical page number** is **hashed** into a **page table**.
- ▶ This **page table** contains a **chain** of elements **hashing to the same location**.
- ▶ Each element contains
 1. The **logical page number**
 2. The value of the **mapped page frame**
 3. A **pointer** to the next element

Hierarchical Paging



Hierarchical Page Tables

- ▶ Break up the **logical address** space into **multiple page tables**.



Hierarchical Page Tables

- ▶ Break up the **logical address** space into **multiple page tables**.
- ▶ A simple technique is a **two-level page table**.



Hierarchical Page Tables

- ▶ Break up the **logical address** space into **multiple page tables**.
- ▶ A simple technique is a **two-level page table**.
- ▶ We then **page** the **page table**.



Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:



Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.



Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.

- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.
- ▶ Thus, a logical address is:

page number		page offset
p_1	p_2	d
12	10	10

Two-Level Paging Example

- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

- ▶ Thus, a logical address is:

page number		page offset
p_1	p_2	d
12	10	10

- ▶ p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.

Two-Level Paging Example

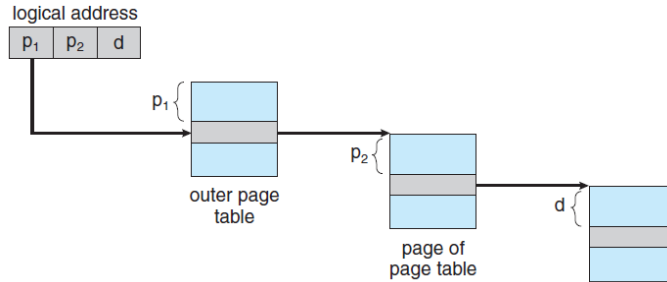
- ▶ A logical address, on 32-bit machine with 1K page size, is divided:
 - A page number consisting of 22 bits.
 - A page offset consisting of 10 bits.
- ▶ Since the page table is paged, the page number is divided into:
 - A 12-bit page number.
 - A 10-bit page offset.

- ▶ Thus, a logical address is:

page number		page offset
p_1	p_2	d
12	10	10

- ▶ p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.
- ▶ Known as forward-mapped page table.

Address-Translation Scheme



Summary



Summary

- ▶ Paging vs. Segmentation



Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages



Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses



Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses
- ▶ Translation Look-aside Buffer (TLB)



Summary

- ▶ Paging vs. Segmentation
- ▶ Physical memory: frames, Logical memory: pages
- ▶ Page table: translates logical to physical addresses
- ▶ Translation Look-aside Buffer (TLB)
- ▶ Page table structure: hierarchical paging, hashed page tables

Questions?

Acknowledgements

Some slides were derived from Avi Silberschatz slides.